

Facebook-Project

October 23, 2016

1 Facebook Check-Ins

FINAL LEADERBOARD SCORE: 0.53454, Team Name: Jong Lee

```
In [21]: import numpy as np
import pandas as pd
from pandas import DataFrame, Series
from sklearn import neighbors, datasets
from sklearn.preprocessing import LabelEncoder
from sklearn.neighbors import KNeighborsClassifier as KNN
```

```
In [34]: #Read data
train = pd.read_csv('./data/train.csv')
test = pd.read_csv('./data/test.csv')
```

```
In [23]: #Take a look at the datasets
print("Training data size is" + str(train.shape))
print("Testing data size is" + str(test.shape))

print(train['accuracy'].describe())
print(train.head(3))
print(test.head(3))
```

Training data size is(29118021, 6)

Testing data size is(8607230, 5)

```
count    2.911802e+07
mean      8.284912e+01
std       1.147518e+02
min       1.000000e+00
25%       2.700000e+01
50%       6.200000e+01
75%       7.500000e+01
max       1.033000e+03
```

Name: accuracy, dtype: float64

	row_id	x	y	accuracy	time	place_id
0	0	0.7941	9.0809	54	470702	8523065625
1	1	5.9567	4.7968	13	186555	1757726713

2	2	8.3078	7.0407	74	322648	1137537235
	row_id	x	y	accuracy	time	
0	0	0.1675	1.3608	107	930883	
1	1	7.3909	2.5301	35	893017	
2	2	8.0978	2.3473	62	976933	

```
In [24]: train.describe()
```

```
Out [24]:
```

	row_id	x	y	accuracy	time
count	2.911802e+07	2.911802e+07	2.911802e+07	2.911802e+07	2.911802e+07
mean	1.455901e+07	4.999770e+00	5.001814e+00	8.284912e+01	4.170104e+01
std	8.405649e+06	2.857601e+00	2.887505e+00	1.147518e+02	2.311761e+01
min	0.000000e+00	0.000000e+00	0.000000e+00	1.000000e+00	1.000000e+00
25%	7.279505e+06	2.534700e+00	2.496700e+00	2.700000e+01	2.030570e+01
50%	1.455901e+07	5.009100e+00	4.988300e+00	6.200000e+01	4.339220e+01
75%	2.183852e+07	7.461400e+00	7.510300e+00	7.500000e+01	6.204910e+01
max	2.911802e+07	1.000000e+01	1.000000e+01	1.033000e+03	7.862390e+01

	place_id
count	2.911802e+07
mean	5.493787e+09
std	2.611088e+09
min	1.000016e+09
25%	3.222911e+09
50%	5.518573e+09
75%	7.764307e+09
max	9.999932e+09

```
In [25]: #Separate spatial data into grid for computation in smaller bits
         #Code mostly taken from 'Sandro' on Kaggle Kernels
```

```
def prepare_data(df, n_cell_x, n_cell_y):
    # """
    #     Feature engineering and computation of the grid.
    #     """
    #Creating the grid
    size_x = 10. / n_cell_x #divide 10 (x values in [0,10]) by number of cells
    size_y = 10. / n_cell_y #divide by 10 (y values in [0, 10]) by number of cells
    eps = 0.00001 #why do eps? Just so for values less than this, just do not
    xs = np.where(df.x.values < eps, 0, df.x.values - eps)
    ys = np.where(df.y.values < eps, 0, df.y.values - eps)
    pos_x = (xs / size_x).astype(np.int64) #changing position for x into number of cells
    pos_y = (ys / size_y).astype(np.int64) #changing position for y into number of cells
    df['grid_cell'] = (pos_y * n_cell_x + pos_x) #the grid # (how many y's in the x's cell)

    #Feature engineering
    fw = [500, 500, 4, 3, 1./22., 2, 10] #feature weights (black magic here)
```

```

#each fw represents x, y, hour, weekday, day, month, year changing
#Note: Changed y fw from 1000 to 750 b/c exceeds int64 and produce
#thing to keep in mind: n_cell_x can't be too big, then will have

```

```

df.x = df.x.values * fw[0]
df.y = df.y.values * fw[1]
initial_date = np.datetime64('2014-01-01T01:01', dtype='datetime64[m]')
d_times = pd.DatetimeIndex(initial_date + np.timedelta64(int(mn), 'm')
                           for mn in df.time.values)
df['hour'] = d_times.hour * fw[2]
df['weekday'] = d_times.weekday * fw[3]
df['day'] = (d_times.dayofyear * fw[4]).astype(int)
df['month'] = d_times.month * fw[5]
df['year'] = (d_times.year - 2013) * fw[6]

df = df.drop(['time'], axis=1) #drop time b/c converted into time
return df

```

```

In [26]: #Also code adapted from 'Sandro'
        #jk most of this stuff is mine now

```

```

def process_one_cell(df_train, df_test, grid_id, threshold, n_cell_x, n_cell_y):
    """
    Throw in a training dataset and it will split it into local training and test sets
    do a KNN classification inside one grid cell.
    """

    #Working on df_train to train ONE CELL (from grid_cell column)
    df_cell_train = df_train.loc[df_train.grid_cell == grid_id] #gets all data for that cell
    place_counts = df_cell_train.place_id.value_counts()
    #getting counts of places in that cell

    mask = (place_counts[df_cell_train.place_id.values] >= threshold).values
    df_cell_train = df_cell_train.loc[mask] #weeds out ID's with less than threshold

    #Working on df_test
    df_cell_test = df_test.loc[df_test.grid_cell == grid_id]

    #Saving row ids of test for our output later, which will be along with accuracy
    row_ids = df_cell_test.index

    features = ['x', 'y', 'hour', 'day', 'weekday', 'month', 'year', 'accuracy']

    train_y = df_cell_train['place_id']
    train_x = df_cell_train[features]

    test_x = df_cell_test[features]

```

```

# KNN algorithm and test accuracy
knn = KNN(15) #15 nearest neighbors
knn.fit(train_x, train_y) #classifying data based on 15-nearest neighbors
all_preds = knn.predict_proba(test_x)

#Saving predictions into preds_per_cell
preds_per_cell = np.zeros((test_x.shape[0], 3), dtype=int)
for record in range(len(all_preds)):
    top3_idx = all_preds[record].argsort()[-3:][::-1]
    preds = knn.classes_[top3_idx]
    preds_per_cell[record] = preds

train_acc = knn.score(train_x, train_y) # score KNN on train set

return preds_per_cell, row_ids, train_acc

In [27]: def process_grid(df_train, df_test, threshold, n_cells, n_cell_x, n_cell_y)
        """
        Iterates over all grid cells, return average training and testing accuracy
        """
        preds = np.zeros((df_test.shape[0], 3), dtype=int)
        small_train_acc_sum = 0
        #small_test_acc_sum = 0

        for grid_id in range(n_cells):
            if grid_id % 100 == 0:
                print('iter: %s' %(grid_id)) #Print iteration per 100 grids
                print(small_train_acc_sum / (grid_id - 1))

                #Applying classifier to one grid cell
                pred_labels, row_ids, small_train_acc = process_one_cell(df_train,
                                                                           grid_id,
                                                                           n_cell_x,
                                                                           n_cell_y)

                small_train_acc_sum += small_train_acc #add up each training accuracy
                #small_test_acc_sum += small_test_acc #add up each testing accuracy

                #Updating predictions
                preds[row_ids] = pred_labels

        train_acc_avg = small_train_acc_sum/n_cells
        #test_acc_avg = small_test_acc_sum/n_cells

        print('Generating submission file ...')

        #Auxiliary dataframe with the 3 best predictions for each sample

```

```

df_aux = pd.DataFrame(preds, dtype=str, columns = ['l1', 'l2', 'l3'])

#Concatenating the 3 predictions for each sample
ds_sub = df_aux.l1.str.cat([df_aux.l2, df_aux.l3], sep=' ')

#Writing to csv
ds_sub.name = 'place_id'
ds_sub.to_csv('sub_knn.csv', index=True, header=True, index_label='row')

return train_acc_avg#, test_acc_avg

```

1.1 Notes

All you need to call is the `process_grid()` method with appropriate parameters

- Inside `process_grid()`, the `process_one_cell()` method is used on each cell, and
- Inside each `process_one_cell()` call, `prepare_data()` is used

1.2 Implementation

```

In [30]: #Required Variables
        n_cell_x = 30
        n_cell_y = 30
        threshold = 3

In [35]: #First, feature engineering + separating training data into amplified grid
        df_train = prepare_data(train, n_cell_x, n_cell_y)
        df_test = prepare_data(test, n_cell_x, n_cell_y)

In [36]: # #Then running classification model through process_grid
        process_grid(df_train, df_test, threshold, n_cell_x * n_cell_y, n_cell_x,

iter: 0
-0.0
iter: 100
0.589275893058
iter: 200
0.583516438543
iter: 300
0.583023747064
iter: 400
0.583135357303
iter: 500
0.582176202611
iter: 600
0.582177203954
iter: 700
0.581937037505

```

```
iter: 800  
0.581684408644  
Generating submission file ...
```

```
Out[36]: 0.58086004301081262
```

```
In [ ]:
```

```
In [ ]:
```